

# Import necessary libraries

```
In [19]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

## Load your dataset

## Replace 'your\_dataset.csv' with the actual path to your dataset

```
In [20]: data = pd.read_csv("predicting_house_prices.csv")
data
```

```
Out[20]:
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.45857	5.682861	7.009188	4.09	23086.80050	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.64245	6.002900	6.730821	3.09	40173.07217	1.505891e+06	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	61287.06718	5.865890	8.512727	5.13	36882.15940	1.058988e+06	9127 Elizabeth Stravenue\nDanieltown, WI 06482...
3	63345.24005	7.188236	5.586729	3.26	34310.24283	1.260617e+06	USS Barnett\nFPO AP 44820
4	59982.19723	5.040555	7.839388	4.23	26354.10947	6.309435e+05	USNS Raymond\nFPO AE 09386
...	...	...	...	...	...	...	...
4995	60567.94414	7.830362	6.137356	3.46	22837.36103	1.060194e+06	USNS Williams\nFPO AP 30153-7653
4996	78491.27543	6.999135	6.576763	4.02	25616.11549	1.482618e+06	PSC 9258, Box 8489\nAPO AA 42991-3352
4997	63390.68689	7.250591	4.805081	2.13	33266.14549	1.030730e+06	4215 Tracy Garden Suite 076\nJoshualand, VA 01...
4998	68001.33124	5.534388	7.130144	5.44	42625.62016	1.198657e+06	USS Wallace\nFPO AE 73316
4999	65510.58180	5.992305	6.792336	4.07	46501.28380	1.298950e+06	37778 George Ridges Apt. 509\nEast Holly, NV 2...

5000 rows × 7 columns

## Define the features (independent variables) and the target variable (house prices)

```
In [24]: X = data[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms', 'Avg. Area Number of Bedrooms']
y = data['Price']
y
```

```
Out[24]:
```

0	1.059034e+06
1	1.505891e+06
2	1.058988e+06
3	1.260617e+06
4	6.309435e+05
...	...
4995	1.060194e+06
4996	1.482618e+06
4997	1.030730e+06
4998	1.198657e+06
4999	1.298950e+06

Name: Price, Length: 5000, dtype: float64

## Split the data into training and testing sets

```
In [25]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

## Create and train a Linear Regression model

```
In [26]: model = LinearRegression()
```

```
In [25]: model = LinearRegression()  
model.fit(X_train, y_train)
```

```
Out[26]: ▼ LinearRegression  
LinearRegression()
```

## Make predictions

```
In [27]: y_pred = model.predict(X_test)
```

## Evaluate the model

```
In [29]: mse = mean_squared_error(y_test, y_pred)  
mae = mean_absolute_error(y_test, y_pred)  
r2 = r2_score(y_test, y_pred)
```

## Print the evaluation metrics

```
In [30]: print("Mean Squared Error:", mse)  
print("Mean Absolute Error:", mae)  
print("R-squared:", r2)
```

```
Mean Squared Error: 10089009299.499422  
Mean Absolute Error: 80879.09722218302  
R-squared: 0.917997170698532
```

## You can also examine the model coefficients to see feature importance

```
In [31]: coefficients = model.coef_  
intercept = model.intercept_  
print("Coefficients:", coefficients)  
print("Intercept:", intercept)
```

```
Coefficients: [2.16522058e+01 1.64666481e+05 1.19624012e+05 2.44037761e+03  
1.52703134e+01]  
Intercept: -2635072.90091678
```

```
In [ ]:
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js